

## Heaps and priority queues

As mentioned earlier, an array  $\underline{t}$  (of objects or pointers to objects) can be seen as a rooted binary tree as follows:

- The root node is  $\underline{t}[0]$ .
- The child nodes of  $\underline{t}[i]$  are the nodes  $\underline{t}[2 \cdot i + 1]$  and  $\underline{t}[2 \cdot i + 2]$ .
- The parent node of  $\underline{t}[j]$  is the node  $\underline{t}[i]$  where  $i = \lfloor (j - 1)/2 \rfloor$ .

As usual, the actual storage allocated for the array will be greater than the size of the tree so that there is room for the tree to grow. As a result, we have an index  $\underline{t}.\text{size}$  so that all nodes correspond to indices which are between 0 and  $\underline{t}.\text{size} - 1$ .

A max-heap is a binary tree as above where:

- The nodes can be compared (they lie in an ordered set  $A$ ).
- The child nodes are  $\leq$  the parent node.

In particular, the root node is maximal.

One can similarly talk about a min-heap.

## Operations

The common operations on a max-heap are:

- Given an element  $a$  of  $A$  we have  $\text{push}(\underline{t}, a)$  which inserts  $a$  in a new node position so chosen that the resulting tree is *also* a max-heap.
- We can delete the root node of a max-heap with  $\text{pop}(\underline{t})$  to create a max-heap with one less node. This data at this node is the “return value” of this algorithm.

Both these operations are written in terms of other operations as explained below.

**Possibly damaged max-heap.** The max-heap property at the node  $p$  means the following. Let  $\gamma$  denote the parent of  $p$  (assuming  $p$  is not the root node). Let  $\lambda$  and  $\rho$  be its child nodes (assuming  $p$  is not a leaf node). The max-heap property at  $p$  is the requirement

$$\gamma \geq p \text{ and } p \geq \max\{\lambda, \rho\}$$

We can say that the max-heap property is “possibly up-damaged” at the node  $p$  if we don’t know the relation between  $\gamma$  and  $p$  but  $p \geq \max\{\lambda, \rho\}$  (or  $p$  is a leaf node).

We can say that the max-heap property is “possibly down-damaged” at the node  $p$  if we don’t know the relation between  $p$  and  $\max\{\lambda, \rho\}$  but  $\gamma \geq p$  (or  $p$  is the root node).

The primary operations will fix possible damage at *exactly* one node.

- The operation `siftup( $\underline{t}, p$ )` operates on a tree  $\underline{t}$  for which the max-heap property holds except that it is possibly up-damaged at  $p$ .
- The operation `siftdown ( $\underline{t}, p$ )` operates on a tree  $\underline{t}$  for which the max-heap property holds except that it is possibly down-damaged at  $p$ .

In both cases, the result of these operations is a tree that satisfies the max-heap property at all nodes.

Given these algorithms, we can make `push` by adding the node to be inserted in the left-most unoccupied slot available for new leaf modes. This will create a tree which is possibly up-damaged at this leaf node. So after this the `siftup` operation will produce a max-heap as required.

Similarly, we can make `pop` by replacing the root node by the right-most leaf node and deleting the right-most leaf node. This will create a tree which is possibly down-damaged at the root node. So after this the `siftdown` operation will produce a max-heap as required. Note that `pop`-ping the right-most leaf node is essentially the same as deleting it.

More generally, we can `change` the data  $a$  at a specific node  $p$  to data  $b$  as follows. We first, compare  $a$  with  $b$ . If  $a$  is greater than  $b$ , then replacing  $a$  by  $b$  creates a tree that is down-damaged at  $p$ ; so we only need to `siftdown` to fix it. Similarly, if the value  $b$  is greater than  $a$ , then replace  $a$  by  $b$  creates a tree that is up-damaged at  $p$ ; so we only need to `siftup` to fix it.

We can also `delete` the node  $p$  as follows. We first replace the  $p$ -th node with the right-most leaf node and remove the latter node. We then apply `siftdown` or `siftup` at the  $p$ -th node depending on the comparison of the current node with the deleted node like we did for `change`. In other words, this operation can be considered as follows:

- Delete the right-most leaf node after saving its value as  $b$ .
- Use the `change` operation to replace the  $p$ -th node by  $b$ .

## Algorithms

In the following algorithms we assume that a max-heap (or a binary tree) is implemented as an (expandable) array  $\underline{t} = [t_0, \dots, t_{n-1}]$  of elements that can be compared. (In Python, we can use a `List` data type to represent such an array.)

```
define siftup( $\underline{t}, p$ )
  Set  $\gamma$  as  $\lfloor (p - 1)/2 \rfloor$ .
  if  $\gamma \geq 0$  and  $t_\gamma < t_p$ 
    Swap  $t_p$  and  $t_\gamma$ .
  return siftup( $\underline{t}, \gamma$ )
return  $\underline{t}$ 
```

In this the algorithm  $p$  is the index of the node at which the tree may be up-damaged and  $\gamma$  is the index of its parent node. Note that the possible up-damage is propagated from  $p$  to  $\gamma$  while removing the damage at  $p$ .

The up-damage in the max-heap ascends along branch containing  $p$ . As a result the total number of comparisons and swaps is bounded by the depth  $\lceil \log_2(n) \rceil$  of the tree.

(Note that in Python, if  $\underline{t}$  is implemented as a list, it will be modified in-place. So the **return** is not required.)

```
define siftdown( $t, p$ )
  Set  $n$  to be size of  $\underline{t}$ .
  Set  $q$  as  $p$ .
  Set  $\lambda$  to be  $2 \cdot p + 1$ .
  Set  $\rho$  to be  $2 \cdot p + 2$ .
  if  $\lambda < n$  and  $t_\lambda > t_q$ 
    Set  $q$  to be  $\lambda$ .
  if  $\rho < n$  and  $t_\rho > t_q$ 
    Set  $q$  to be  $\rho$ .
  if  $p$  is different from  $q$ 
    Swap  $a_p$  and  $a_q$ .
  return siftdown( $\underline{t}, q$ )
return  $\underline{t}$ 
```

Here  $p$  represents the node where the possible down-damage is. It's child nodes are  $\lambda$  and  $\rho$ . Next, note that only *one* of the nodes  $\lambda$  or  $\rho$  is (possibly) swapped with  $p$ . The possible down-damage is propagated from  $p$  to the node with which it is swapped while removing the damage at  $p$ .

There are two comparisons of nodes of  $t$  and at most one swap of notes of  $t$  in each call. The possible down-damage in the max-heap descends along *only* one branch. As a result the total number of comparisons is bounded by  $2\lceil \log_2(n) \rceil$  and the number of swaps is bounded by  $\lceil \log_2(n) \rceil$ .

**Exercise:** Modify the pseudo-code above to use **while** constructions instead of **if** and recursion.

As mentioned above, the **push** and **pop** operations can be made using these two algorithms.

```
define push( $t, a$ )
  Set  $n$  to be size of  $\underline{t}$ .
  Append  $a$  to  $\underline{t}$ .
  Call siftup( $\underline{t}, n$ ).

define pop( $\underline{t}$ )
  Set  $n$  to be size of  $\underline{t}$ .
  Save the return value  $a$  which is set as the root element  $t_0$ .
  Set  $t_0$  to be the last leaf node  $t_{n-1}$ .
  Delete the last leaf node from  $\underline{t}$ .
  Call siftdown( $\underline{t}, 0$ ).
return  $a$ 
```

Note that `pop` only works if the tree has at least one node!

More generally, we can `change` the value at the node at index  $p$  in the tree  $\underline{t}$ , or `extract` it.

```
define change( $\underline{t}, p, b$ )
  Save the current value  $a$  of  $t_p$ .
  Set  $t_p$  to be  $b$ .
  if  $t_p > a$ 
    Call siftup( $\underline{t}, p$ ).
  else
    Call siftdown( $\underline{t}, p$ ).

define extract( $\underline{t}, p$ )
  Set  $n$  to be size of  $\underline{t}$ .
  Save the return value  $a$  which is set as  $t_p$ .
  Set  $t_p$  to be the last leaf node  $t_{n-1}$ .
  Delete the last leaf node from  $\underline{t}$ .
  if  $p$  is  $n - 1$ 
    return  $a$ 
  if  $t_p > a$ 
    Call siftup( $\underline{t}, 0$ ).
  else
    Call siftdown( $\underline{t}, p$ ).
  return  $a$ 
```

## Build

Naively, one may build a heap out of an array of data incrementally, using `push`. This will take  $2 \sum_{k=1}^n \lceil \log_2(k) \rceil$  or  $O(n \log(n))$  comparisons. However, there is a better approach to `heapify` an array as follows.

Given an array  $\underline{t}$  of elements of an ordered set  $A$ , we can convert it into a max-heap as follows. Starting at the right-most and deepest non-leaf node, we work our way to the left and upwards in the tree “fixing” the tree that has the given node as root using `siftdown`.

```
define heapify( $\underline{t}$ )
  Set  $n$  as the size of  $\underline{t}$ .
  for  $i$  from  $\lfloor n/2 \rfloor$  decreasing to 0
    Call siftdown( $\underline{t}, i$ ).
```

We note that the call to `siftdown` takes  $2h_i$  steps where  $h_i$  is the depth of the sub-tree rooted at  $i$ . When  $i$  in one of the  $2^k$  indices the in range  $[2^k - 1, 2^{k+1} - 2]$  we have  $h_i = h = \lceil \log_2(n) \rceil - k$ . In other words,  $2^k$  is roughly  $n/2^h$ . Thus, we

see that the number of steps is (roughly)

$$\sum_{h=1}^{n/2} \frac{n}{2^h} 2h = 2n \sum_{h=1}^{n/2} \frac{h}{2^h} \leq 2n \sum_{h=1}^{\infty} \frac{h}{2^h} = 4m$$

**Exercise:** Make the above calculation more precise and calculate the precise number of steps in the worst case.

### Priority Queues

One can use a max-heap to implement queues. Each node is a pair-record consisting of the queued object and a “priority”. Comparison for the max-heap property is based only on the priority.

- `enqueue` is implemented by using `push` to insert a new node whose priority is 1 less than the priority of the last object queued.
- `dequeue` is implemented easily using `pop`.

In fact, since we can choose the priority at the time of insertion *and* change the priority, what we have is a “priority queue” data structure. This has the operations:

- `insert` which inserts an object into the queue with a given priority.
- `maximum` which finds the element with the highest priority.
- `extract` which extracts the element with the highest priority and removes it from the queue.
- `change` which changes the priority of a chosen element of the queue.

From the above discussion it is clear how such a data structure can be implemented using max-heap.